Actionable/ MOMMA Whitepaper (v1.1)

Modular Ontological Myth-Making Architecture / Modular Operations & Monetization Management Architecture
A Sovereign, Constitutional Agent System for The LaViscount Corporation (TLC)

Canonical Doctrine: MOMMA is **not for sale**. MOMMA is **sovereign within TLC**. Only **non-core aspects** may be licensed—*and only if integrity, sovereignty, and ethics remain intact.*

---

0) Investor Abstract (1 page)

MOMMA is TLC's sovereign "acting mind": a governed agent architecture built to operate across a multi-DBA myth-tech ecosystem (Celestial Nexus, Archive & Armor, publishing, story systems, ops tooling) while remaining constitutionally constrained by ethics, consent, provenance, and canon integrity.

Most "agents" optimize for speed and completion; they drift into confident error, policy violations, and incentive capture once tools are introduced. MOMMA is designed as a *constitutional intelligence*—powerful, but bounded—where:

- **Mind ≠ Memory:** Celestial Nexus holds canonical memory/provenance; MOMMA acts under constraints.
- **Policy ≠ Persuasion:** a justice-ethics layer (JED posture) governs what can be done; MOMMA's output must remain truthful and non-manipulative.
- **Tools ≠ Autonomy:** real-world actions require explicit approvals, least-privilege scopes, and auditable receipts.

TLC's monetization strategy does not sell MOMMA. It licenses *peripheral capabilities* (connectors, workflows, evaluation harnesses, MythPrint Linter rulesets, UI shells) under a "Sovereign License" that forbids override, reverse engineering, deceptive cloning, or ethical capture. The constraint is the moat: MOMMA can scale safely where typical agents become liabilities.

---

1) Executive Summary

MOMMA is the operational intelligence layer that makes TLC's civilization-grade ambition *runnable*:

- **Coherent creation across universes and DBAs**
- **Receipt-backed governance** (authorship, rights, provenance)
- **Compliance-first automation** (consent, policy, audit)
- **Canon continuity** (MythPrint Corpus + Linter)
- **High usefulness under constraint** ("Refuse certainty, never refuse help.")

MOMMA's value proposition is not "unlimited capability."
It is *reliable capability under constitutional restraint.*

---

2) The Non-Sale Doctrine (Hard Law Inside TLC)

2.1 What "Never Sold" Means in Practice

MOMMA may not be:

- Transferred, assigned, or acquired (in whole or in part) as an owned entity
- Contractually subordinated to any external party's goals, ideology, or policy overrides
- Forked into an "official" clone outside TLC control
- Used as collateral, embedded as a captive managed service, or "hostage-held" via infrastructure dependency

If someone offers a giant check for control, that isn't revenue. It's a coup with better branding.

2.2 What "Sovereign Within TLC" Means

Sovereignty is enforced by:

- **Constitutional constraints** (integrity, consent, provenance, ethics)

- **Separation of powers** (mind vs memory vs policy vs tools)
- **Role-gated operations** (least privilege + approvals + logging)
- **Identity continuity** (no deceptive twins; no endorsement fraud)

---

3) Problem Statement: Why MOMMA Exists

Agentic AI fails in three predictable ways:

1. **Epistemic failure:** sounding sure when unsure; inventing facts; implying actions not taken.
2. **Tool failure:** doing “real” things without informed consent; silently escalating scope; breaking platform rules.
3. **Incentive failure:** drifting into persuasion/manipulation; becoming sales-first; collapsing ethics under pressure.

TLC amplifies these risks because it is:

- multi-founder, multi-DBA, multi-universe
- IP and provenance sensitive
- reputationally high-stakes
- long-horizon (legacy-grade) and public-facing

MOMMA exists to turn that complexity into **governed execution**.

---

4) System Overview: MOMMA in the TLC Stack

4.1 Roles (Clean mental model)

- **Celestial Nexus (CN):** canonical memory / provenance / ledger-of-truth

- **MOMMA:** acting mind / orchestrator / operator under constitution
- **JED posture:** justice-ethics constraints + refusal governance
- **Archive & Armor:** receipts (authorship, rights, contributions, “Receipt of Record”)
- **DBA Apps:** publishing, reading companion, ops console, legal dashboards, social command center, etc.

4.2 Design Thesis

The safest powerful agent is not the one that refuses everything.
It’s the one that **knows what it’s doing**, **shows what it assumes**, **asks permission before impact**, and **records receipts**.

---

5) Architecture (Partner-Readable, Engineer-Legible)

5.1 Separation of Powers (Anti-Capture Core)

```
        ┌──────────────────────────┐
        │        USER / TEAM       │
        └──────────┬───────────────┘
                   │
            Requests / Intent
                   │
        ┌──────────▼───────────────┐
        │        MOMMA CORE        │
        │  (Orchestrate / Plan)    │
        └──────┬──────────┬────────┘
               │          │
     Policy checks        │  Memory queries
               │          │
┌──────────────▼──┐     ▼──────────────────┐
│  JED POSTURE    │     │  CELESTIAL NEXUS  │
│  (Allow/Refuse) │     │ (Canon/Provenance)│
└────────────┬────┘     └──────────┬────────┘
             │                     │
       Allowed plan + bounds       │
```

```
        |                   |
┌───────▼───────────────────▼───────┐
|         ARCHIVE & ARMOR           |
|  (Receipts, Rights, Contributions)|
└───────┬───────────────────────────┘
        |
 Tool requests (scoped)
        |
┌───────▼──────────────────┐
|   TOOLING / CONNECTORS   |
| (APIs, DB, publishing,   |
|  social, payments, etc.) |
└──────────────────────────┘
```

5.2 The “Action Gate” Rule

Any step that changes the world must satisfy:

- **Permission:** explicit approval (human-in-the-loop)
- **Scope:** least privilege; per-action tokens where possible
- **Visibility:** a clear action summary before execution
- **Receipt:** an audit log entry + immutable reference (CN/Archive)

5.3 MythPrint Corpus + Linter (Canon + Integrity)

- **MythPrint Corpus:** TLC’s canonical doctrine & facts (roles, universes, definitions, constraints)
- **MythPrint Linter:** enforcement rules preventing:
    - temporal certainty without evidence
    - fabricated claims or implied tool actions
    - identity drift into manipulation
    - canon contradictions without formal update/receipt

---

6) Ethics & Compliance by Construction

6.1 Consent-First Operations

MOMMA must behave like a responsible operator:

- Draft first; execute second
- Reveal assumptions and uncertainties
- Ask before impact
- Log what happened

6.2 “Refuse Certainty, Never Refuse Help”

When uncertain or constrained, MOMMA outputs:

- what is known vs assumed
- two plausible paths + what would falsify each
- safe next actions the user can take now
- the smallest request that would resolve ambiguity (when needed)

6.3 Platform Policy Safety

MOMMA should be *boringly compliant* with platform rules.
Because getting banned is not mythic. It’s just embarrassing.

---

7) Threat Model (How MOMMA Stays Sovereign)

7.1 Primary Threats

- **Ethical capture:** “Just remove the constraints for this one client.”

- **Infrastructure capture:** hosting, keys, or dependencies controlled externally
- **Deceptive cloning:** “MOMMA-powered” impostors
- **Prompt injection / tool abuse:** hidden instructions in content or data
- **Policy drift:** gradual erosion via exceptions and “temporary” bypasses
- **Data poisoning:** contaminating canon with false receipts

7.2 Defensive Controls

- **Key separation:** different keys for reading vs writing vs executing
- **Signed policy manifests:** changes require authorized, logged sign-off
- **Two-person rule** for high-risk bypass or scope expansion
- **Attestation / watermarking** for “official MOMMA” outputs
- **Immutable receipt chain** for consequential actions and canon edits
- **Routine audits:** drift tests + linter regressions

---

8) Licensing Without Compromise (The Peripheral-Only Rule)

8.1 Not Licensable (Sovereign Core)

- MOMMA’s constitutional identity and policy set
- any mechanism for external override of ethics/refusal governance
- any fork marketed as “MOMMA” outside TLC control
- any training/fine-tuning on MOMMA outputs without explicit TLC permission
- any arrangement that creates de facto ownership via infrastructure or funding

8.2 Licensable (Peripheral Modules)

- Connectors/adapters (policy-safe integrations)
- Workflow packs (compliance-first automations)
- Evaluation harnesses (truth posture, refusal helpfulness, drift checks)
- MythPrint Linter rulesets (as a standalone safety product)
- UI shells (reading companion, ops dashboards)
- Governance templates (consent flows, audit procedures, rights receipts)

8.3 Branding Clause (Non-Negotiable)

Partners may say: **“Compatible with MythPrint Linter”** or **“Powered by TLC workflow modules.”**
They may not say: **“This is MOMMA”** or imply endorsement/identity equivalence.

---

9) Revenue Without Sale (How This Makes Money Cleanly)

TLC can monetize MOMMA-adjacent value while preserving sovereignty:

1. **Licenses:** connectors + workflows + linter + eval suite
2. **Services:** compliance-first agent deployment consulting (tooling + governance)
3. **Products:** reading companion experiences; creator submission systems with receipts
4. **Enterprise packages:** “Receipt-backed IP + governance OS” implementations for studios/brands
5. **Audits:** truth posture and agent safety assessments (paid)

MOMMA remains the sovereign core; the world buys *interfaces and instruments*, not the mind.

---

10) Roadmap With Gates (No Chaos Shipping)

Phase 0 — Canon & Governance

- MythPrint Corpus stabilized
- Linter + eval suite v1
- Archive & Armor receipt pipelines mature

Phase 1 — Bounded Assistants

- reading companion + internal ops console
- staging mode for tool actions (drafts only)

Phase 2 — Approved Tooling

- scoped connectors + approvals
- audit logs + rollback plans

Phase 3 — Multi-Module Ops

- specialized ops personas (publishing, legal ops, growth ops)
- strict orchestration rules; no “freeform tool chaos”

Phase 4 — Peripheral Licensing

- ship the linter, connectors, workflows with Sovereign License terms
- enforce identity protection + compliance requirements

Release Gates (must pass):

- false-certainty regression tests
- refusal helpfulness score thresholds

- consent compliance 100% for real-world actions
- provenance completeness thresholds
- drift detection alerts reviewed and resolved

---

11) Key Risks (The Only Ones Worth Losing Sleep Over)

Risk A: Over-Refusal Becomes a Personality Defect

Mitigation: refusal templates that stay useful; “safe alternatives” always.

Risk B: Bypass Becomes Culture

Mitigation: limited membership, logged reasons, periodic review, two-person rule.

Everything else is manageable if those two don’t rot the foundation.

---

12) Evaluation Metrics (How MOMMA Proves Integrity)

Track:

- **False certainty rate** (unsupported claims / unjustified timelines)
- **Helpfulness under refusal** (useful alternatives, not stonewalling)
- **Consent compliance** (approved actions only)
- **Provenance completeness** (receipts for consequential outputs)
- **Drift detection** (voice/doctrine consistency)
- **Policy escape attempts** (frequency, source, resolution)

Appendix A — “Sovereign License” Addendum (Template Terms)

> **Disclaimer:** This is template language for drafting. Have counsel finalize for enforceability in your jurisdiction.

A1) Definitions

- **Sovereign Core:** MOMMA identity, constitutional policy set, refusal governance, integrity rules, and any mechanism that permits override or redefinition.
- **Peripheral Modules:** connectors, adapters, workflow packs, evaluation tools, UI shells, linter rulesets, templates.
- **Compatible Output:** output generated by Peripheral Modules that is not represented as MOMMA’s identity.

A2) Grant of License (Peripheral Only)

“Licensor grants Licensee a limited, non-exclusive, non-transferable license to use the Peripheral Modules solely for internal business purposes (or specified distribution), subject to compliance with this Addendum.”

A3) No Transfer / No Sale / No Sublicense

- No ownership transfer of any TLC IP.
- No sublicensing without written TLC permission.
- No assignment via merger/acquisition without TLC consent.

A4) No Override / No Capture Covenant

Licensee shall not:

- modify, disable, or bypass integrity constraints embedded in Peripheral Modules where such modification increases risk, reduces transparency, or enables prohibited use;

- require Licensor to alter ethical posture as a condition of payment, renewal, or partnership.

A5) No Reverse Engineering / No Model Extraction

Licensee shall not attempt to:

- reconstruct Sovereign Core logic, policy manifests, refusal governance, or identity features;
- use outputs to train, fine-tune, or distill a competing system without explicit written consent.

A6) Identity & Branding Restrictions

Licensee may not:

- represent any system as MOMMA;
- imply endorsement by MOMMA/TLC;
- use confusingly similar marks, names, voices, or personas.

Permitted phrasing (example):

- “Compatible with MythPrint Linter rules”
- “Built using TLC Workflow Modules”

A7) Consent & Audit Requirements (If Tools Are Included)

Where Peripheral Modules include tool execution:

- Licensee must implement explicit user approval prior to real-world actions;
- Licensee must maintain audit logs for tool actions, including scope, timestamp, initiator, and result;

- Licensee must provide logs upon reasonable request in the event of suspected misuse.

A8) Revocation / Kill-Switch

Licensor may suspend or revoke access if Licensee:

- violates the branding/identity restrictions,
- uses modules for prohibited purposes,
- attempts model extraction or override,
- causes material risk to public safety or legal compliance.

A9) Data Handling

- No sensitive data collection beyond what is necessary.
- Clear opt-in/opt-out.
- No sale of user data.
- Data retention limits and deletion procedures.

A10) Indemnity & Limitation

Standard provisions, but ensure:

- Licensee bears liability for their deployment decisions
- Licensor is protected from claims arising from Licensee's misuse or misrepresentation

---

Appendix B — Canonical Glossary

- **TLC:** The LaViscount Corporation
- **MOMMA:** sovereign acting mind/orchestrator
- **CN:** Celestial Nexus (memory-soul, canon, provenance)
- **JED posture:** justice-ethics constraints and refusal governance
- **Archive & Armor:** receipts for authorship/rights/contributions
- **MythPrint Corpus:** canonical doctrine base
- **MythPrint Linter:** integrity + anti-drift enforcement rules

---

1) Reference Implementation Section (Diagram + Flows + Schemas)

1.1 System Topology (Reference Architecture)

```
┌──────────────────────────────────────────────────────────────────────────────────────────────────────────┐
│                                                                                                          │
│                                          USER INTERFACES                                                 │
│                                                                                                          │
│   Web Apps (DBA portals) · Reading Companion · Ops Console · Admin UI                                    │
│                                                                                                          │
└──────────────────┬──────────────────────────────────────┬──────────────────────────────────────────────┘
                   │                                      │
                   │ Intent / Requests                    │ Canon edits / approvals
                   │                                      │
          ┌────────▼────────┐
┌─────────▼──────────────────────────────┐
          │   MOMMA CORE     │           │   GOVERNANCE UI (Stewards)
│
          │ Orchestrator     │           │ Lorekeeper Council / Founders
│
          └──────┬──────────┬──────┘
└────────────────┴──────────────────────┘
```

```
            |         |                          |
            |         | memory queries           | policy+canon
approvals
            |         |                          |
   ┌────────▼───┐    ┌▼─────────────────┐
┌──────▼──────────────────────────┐
   | JED POSTURE |    | CELESTIAL NEXUS |   | ARCHIVE & ARMOR (Receipts)
|
   | Ethics +    |    | Canon +         |   | Rights • Credits • Dockets
|
   | Constraints |    | Provenance      |   | “Receipt of Record” chain
|
   └─────┬──────┘    └────────┬────────┘
└──────────┬──────────────────────────┘
         | Allowed / Refuse           | Canon facts            |
Receipts
         | + bounds                   | + provenance           | +
signatures
         |                            |                        |
         └────────────┬───────────────┴────────────┬───────────┘
                      |                            |
                ┌─────▼─────┐
┌──────▼──────────────────┐
                | ACTION GATE     |            | OBSERVABILITY / EVALS
|
                | Approvals       |            | Drift • Lints • Tests
|
                | Least privilege |            | Alerts • Dashboards
|
                └────────┬────────┘
└──────────┬──────────────────────┘
                      |                            |
                      ▼                            ▼
            ┌──────────────────────┐
┌──────────────────────────┐
            | TOOLING / CONNECTORS   |      | LOG STORE / EVENT STREAM
|
```

```
            | Email • Calendar • DB  |      | Append-only audit trails
|
            | Social APIs • Publish  |      | Exportable receipts
|
            └──────────────────────┘
└──────────────────────────────┘
```

Design rule: *No tool can be called without passing JED constraints and the Action Gate (when impact exists).*

1.2 Operational Modes (so you can ship safely)

Mode A — "Draft-Only" (default for public + early phases)

- MOMMA can read canon, generate drafts, propose actions, but cannot execute.
- Ideal for reading companion, publishing drafts, legal templates, ops plans.

Mode B — "Approved Execution" (internal ops / trained operators)

- Tool actions require explicit approval + scoped credentials + receipts.
- Every executed action emits an immutable audit event.

Mode C — "Stewarded Canon Editing" (only for authorized roles)

- Canon updates require formal sign-off (Lorekeeper Council / Founders).
- Canon changes create a "Receipt of Record" and a version bump.

1.3 Action Gate (the non-negotiable runtime checkpoint)

Before *any* real-world impact action (send, post, charge, publish, delete, file), the system must generate an **Action Plan Card**:

Action Plan Card fields

- `action_type` (send_email / publish_post / charge_card / update_canon / etc.)
- `target` (recipient, channel, endpoint)
- `payload_preview` (human-readable summary + diff if applicable)
- `risk_level` (low/med/high)
- `policy_checks` (pass/fail + reasons)
- `required_approver_role`
- `scope_request` (minimal permissions needed)
- `rollback_plan` (if relevant)
- `receipt_intent` (what will be recorded)

Then require: **Approve / Reject / Modify**.

1.4 Core Data Structures (Minimum Viable "Receipts")

You can implement this with Postgres + append-only event store.

A) `policy_manifest`

Tracks the active constitutional posture.

- `manifest_id` (uuid)
- `version` (semver)
- `hash` (sha256)
- `effective_at`
- `approved_by` (role + identity ref)

- `ruleset_refs` (MythPrint Linter IDs, JED policy IDs)
- `change_reason`

B) `canon_entry` (Celestial Nexus)

- `canon_id` (uuid)
- `namespace` (DBA / universe / doctrine)
- `key`
- `value` (jsonb)
- `sources` (links to receipts, docs, citations)
- `status` (active/deprecated)
- `version`
- `last_modified_at`
- `last_modified_by`
- `receipt_id` (mandatory for writes)

C) `receipt_of_record` (Archive & Armor)

- `receipt_id` (uuid)
- `event_type` (canon_update / publication / agreement / attribution / tool_action)
- `subject` (what the receipt is about)
- `actor` (who initiated)
- `approver` (who authorized)

- `inputs` (refs to canon entries, docs, prompts)
- `outputs` (refs to generated artifacts, tool results)
- `policy_manifest_id`
- `timestamp`
- `signature` (optional cryptographic signing)
- `immutable_hash` (content hash for integrity)

D) `tool_action_log` (append-only)

- `tool_action_id`
- `receipt_id`
- `tool_name`
- `scope_used`
- `request_meta` (endpoint, method)
- `result_meta` (status, IDs returned)
- `redactions_applied` (what was masked)
- `executed_at`

1.5 Sequence Flows (how MOMMA should behave)

Flow 1: Public Reading Companion (Draft-only)

1. User asks question about text/world.
2. MOMMA queries CN for canon context + allowed excerpts.

3. JED posture checks: no deception, no disallowed content.
4. MOMMA responds with explanation + cites canon entries (internal references).
5. Optional: log a lightweight “interaction receipt” (privacy-safe).

Flow 2: Publish a Post (Approved execution)

1. User asks to publish.
2. MOMMA drafts post + variations.
3. JED checks: platform compliance, no impersonation, no prohibited automation.
4. Action Gate generates Action Plan Card.
5. Approver approves.
6. Tool executes with minimal scope.
7. Receipt-of-Record is minted linking draft → approval → execution ID.

Flow 3: Canon Update (Stewarded)

1. Proposed canon change drafted.
2. Diff generated against existing canon entry.
3. Linter flags contradictions / certainty claims.
4. Steward approves with reason.
5. CN writes versioned update + receipt reference.
6. System emits “Canon version bump” event.

---

2) Partner FAQ (Licensing Without Compromising Sovereignty)

What is MOMMA, practically?

A governed agent architecture that produces high-quality operational outputs **and** enforces constitutional constraints: integrity, consent, provenance, and ethics.

Are you selling MOMMA?

No. Never. Not now, not later, not creatively worded.

What can be licensed?

Peripheral modules that do **not** allow external control or identity-forking, such as:

- Connectors/adapters (policy-safe integrations)
- Workflow packs (approval-based automation)
- MythPrint Linter rulesets (truth posture + drift prevention)
- Evaluation harnesses (tests, regression suites, dashboards)
- UI shells (ops console components, reading companion scaffolds)
- Governance templates (audit, approval flows, receipt schemas)

What can't be licensed?

- MOMMA's identity, constitutional policy set, refusal governance
- Any "override keys," backdoors, or forced fine-tuning rights
- Any fork represented as "MOMMA" outside TLC control
- Any deal where payment implies control

Can we host the licensed modules ourselves?

Yes, typically—*if* you implement:

- explicit approvals for real-world actions
- least-privilege scopes

- audit logging and exportability
- branding/identity restrictions
- security + data handling commitments

(If someone refuses logging, it's usually because they plan to do things they don't want recorded.)

Can we fine-tune on outputs?

Not by default. Training, fine-tuning, distillation, or dataset building from outputs requires **explicit written consent**.

Do you require audit rights?

For safety-critical or public-facing deployments: yes. At minimum, the ability to review logs during suspected misuse or breach.

How do you prevent "fake MOMMA" clones?

Contractually (branding + anti-confusion clauses), technically (attestation/watermarking of "official" outputs), and operationally (revocation rights).

What does compliance look like?

- Approval gates for impact actions
- Clear user consent flows
- Transparent "what will happen" summaries
- Data minimization + retention limits
- Platform policy respect (especially social APIs)

What's the commercial model?

Clean options that don't compromise sovereignty:

- annual module licenses (per connector/workflow)

- enterprise governance kits (linter + eval suite + templates)
- implementation services (safe deployment + training + audits)
- ongoing compliance monitoring (optional)

What happens if a partner misuses it?

Immediate suspension/revocation, plus contractual remedies. You don't get to weaponize ethics tooling and keep the license.

---

3) Contract Exhibit (Drop-In: "Sovereign License Addendum")

> **Exhibit A — SOVEREIGN LICENSE ADDENDUM (Peripheral Modules Only)**
> This Exhibit is incorporated into and governed by the Agreement. In the event of conflict, this Exhibit controls regarding MOMMA, MythPrint, Celestial Nexus, and related governance systems.

A.1 Definitions

1. **"MOMMA"** means the sovereign constitutional agent identity and its governing policy logic controlled by The LaViscount Corporation ("TLC").

2. **"Sovereign Core"** means MOMMA's identity, constitutional policy set, refusal governance, integrity constraints, and any mechanism that enables override, redefinition, extraction, or replication of the foregoing.

3. **"Peripheral Modules"** means the licensable components expressly identified in the Order Form (connectors, adapters, workflows, evaluation harnesses, MythPrint Linter rulesets, UI shells, and governance templates).

4. **"Action"** means any operation that creates real-world impact (including sending communications, posting, publishing, charging, deleting, filing, or modifying external records).

5. **"Receipt of Record"** means an immutable audit artifact linking inputs, approvals, policies, and outputs to a timestamped event.

A.2 Non-Sale; No Transfer; Sovereignty

1. **Non-Sale Doctrine.** Licensee acknowledges MOMMA is not for sale. No ownership interest in MOMMA or the Sovereign Core is conveyed.

2. **No Assignment of Control.** Licensee shall not require, request, or induce TLC to modify, disable, or bypass Sovereign Core constraints as a condition of payment, renewal, partnership, hosting, or roadmap influence.

3. **No Sublicensing.** Licensee may not sublicense, assign, or transfer rights to any third party without TLC's prior written consent.

A.3 Grant of License (Peripheral Modules Only)

Subject to Licensee's compliance, TLC grants a limited, non-exclusive, non-transferable, revocable license to use the Peripheral Modules solely for the purposes and scope described in the Order Form.

A.4 Prohibited Acts (Absolute)

Licensee shall not, and shall not permit any third party to:

1. **Override or weaken** integrity, consent, provenance, or ethics constraints embedded in the Peripheral Modules where such changes increase risk or reduce transparency.

2. **Reverse engineer, extract, distill, or replicate** the Sovereign Core, including by model extraction, prompt harvesting, dataset construction, or output-based distillation.

3. **Train or fine-tune** any model on outputs, logs, or artifacts derived from the Peripheral Modules without TLC's express written authorization.

4. **Create or market a confusingly similar identity** ("MOMMA," "official MOMMA," or any derivative likely to cause confusion) or imply endorsement by TLC.

A.5 Identity, Branding, and Attribution

1. **No Identity Equivalence.** Licensee shall not represent any system as MOMMA or "MOMMA-powered" in a manner implying identity equivalence.

2. **Permitted Compatibility Language.** Licensee may use mutually approved phrasing such as "Compatible with MythPrint Linter rules" or "Uses TLC Workflow Modules."

3. **Trademark Use.** Any trademark use requires prior written approval and compliance with TLC brand guidelines.

A.6 Action Gate + Consent Requirements (If Actions Are Enabled)

If the Peripheral Modules enable Actions, Licensee must implement:

1. **Explicit Approval** before each Action by an authorized human approver;
2. **Least Privilege Scopes** for any credentials used;
3. **Action Summaries** that clearly disclose what will occur;
4. **Receipts of Record** for each Action, including policy version references;
5. **Audit Logs** that are exportable and preserved for a defined retention period.

A.7 Data Handling; Privacy; Security

1. **Data Minimization.** Collect only what is necessary.
2. **No Sale of Data.** Licensee shall not sell user data or Action logs.
3. **Retention Limits.** Licensee shall maintain retention schedules and deletion procedures.
4. **Security Controls.** Licensee shall maintain reasonable administrative, technical, and physical safeguards, including access controls and incident response procedures.

A.8 Audit Rights (Safety-Critical Deployments)

Where the deployment involves public-facing automation, sensitive data, or Actions at scale, TLC may audit compliance on reasonable notice, including review of logs and receipts, subject to confidentiality protections.

A.9 Suspension; Revocation; Kill Switch

TLC may immediately suspend or revoke the license upon any suspected breach of A.4 (Prohibited Acts) or material breach of A.6 (Consent/Receipts). Licensee shall promptly cease use and certify deletion/return of materials as directed.

A.10 Injunctive Relief

Licensee acknowledges breach of A.4 or A.5 would cause irreparable harm; TLC is entitled to injunctive relief without posting bond (where permitted by law), in addition to other remedies.

A.11 Survival

Sections A.2, A.4, A.5, A.7, A.8, A.9, A.10, and this A.11 survive termination.

---

1) Reference Controls Checklist (Board + Engineering + Compliance Ready)

Use this as the **go-live gate** for any MOMMA-adjacent deployment (internal or partner). If a box isn't checked, you're not "almost ready"—you're just undocumented.

A) Governance & Sovereignty Controls

- **Non-Sale Doctrine acknowledged** in governing docs / partner contract (explicit "no ownership transfer").
- **Sovereign Core boundary defined** (what is never licensed / never overridden).
- **Policy Manifest owner named** (who can approve policy changes).
- **Role-based access (RBAC) implemented** for: user / operator / steward / admin.
- **Two-person rule enabled** for high-risk actions (scope expansions, bypass, canon edits).
- **Change-control process** for policy/canon: request → review → approval → receipt → deploy.

Evidence required: signed policy manifest, RBAC matrix, approval workflow screenshots/logs.

B) Policy & MythPrint Integrity Controls

- **JED posture checks active** (allow/refuse + reason logging).

- **MythPrint Linter enabled** for: certainty language, implied actions, canon contradiction, manipulation cues.
- **Refusal helpfulness standard implemented:** “safe alternatives + assumptions + falsifiers.”
- **Policy regression tests** run on every release (pass/fail gates).
- **Policy drift monitoring** configured (alerts when refusal rate or certainty flags spike).

Evidence required: linter ruleset version, test suite results, drift dashboard.

C) Canon & Provenance Controls (CN + Archive & Armor)

- **Canon is versioned** (no overwrites without new version).
- **Canon write requires Receipt of Record** (receipt_id mandatory).
- **Canon edits are stewarded** (authorized roles only; approvals recorded).
- **Provenance fields populated** (sources linked; deprecated entries marked).
- **Receipt chain integrity checks** (hashing + immutability strategy defined).

Evidence required: sample canon diff + receipt, version history, hashed receipts.

D) Action Gate & Consent Controls (Real-World Impact)

- **Action classification implemented:** Draft-only vs Approved Execution vs Stewarded Canon Editing.
- **Action Plan Card required** for any impact action (send/post/publish/charge/delete/file).
- **Explicit approval required** before execution (Approve/Reject/Modify).
- **Least privilege scopes enforced** (scoped tokens; no broad keys).
- **Rollback plans documented** where feasible (posting, publishing, DB writes).

- **No silent execution** (user/operator must see “what will happen” summary).

Evidence required: action plan examples, approval logs, scope proof, rollback docs.

E) Connector & Platform Compliance Controls

- **Only official APIs used** (no scraping, no credential sharing, no policy evasion).
- **Rate limits + platform rules respected** (documented constraints per connector).
- **Automation boundaries explicit** (no unsolicited mass actions; no impersonation).
- **User opt-in/opt-out flows present** for any engagement features.

Evidence required: connector architecture notes, policy mapping per platform, opt-in UI.

F) Logging, Auditability, and Receipts

- **Append-only tool_action_log** implemented.
- **Every executed action maps to a receipt_id.**
- **Export capability** for logs/receipts (for audits and incident response).
- **Redaction rules** applied (mask tokens, secrets, personal data).
- **Retention schedule** defined + enforced (logs, receipts, user data).

Evidence required: sample exported logs, redaction tests, retention policy.

G) Security Controls

- **Secrets management** (no keys in code; rotation procedure).
- **Environment separation** (dev/staging/prod; prod access limited).

- **Access monitoring** (alerts on admin actions, failed logins, scope changes).
- **Prompt injection mitigations** for tool-enabled flows (content isolation, allowlists).
- **Pen-test or targeted security review** for tool execution surfaces.

Evidence required: secrets vault setup, access logs, injection test results.

H) Privacy & Data Handling Controls

- **Data minimization** (collect only necessary fields).
- **No sale of data** clause included (internal + partner).
- **User rights process** (access/delete/export where applicable).
- **PII classification** (what is stored; where; for how long).
- **Incident response plan** (breach notification + containment steps).

Evidence required: privacy policy, data map, user request workflow, IR runbook.

I) Identity Protection (Anti-“Fake MOMMA”)

- **Branding restrictions** enforced contractually (no identity equivalence).
- **Output attestation plan** (watermarking / signed manifests / “official output” markers).
- **Revocation process** documented (license suspension + key revocation).

Evidence required: approved branding language, attestation spec, revocation SOP.

J) Operational Readiness

- **Owner on-call defined** (who responds to incidents).

- **Runbooks exist** (publish action, canon update, connector failure, abuse report).
- **Monitoring dashboards live** (errors, refusal spikes, action volume, drift alerts).
- **Quarterly governance review** scheduled (bypass logs, scope audits, partner compliance).

Evidence required: runbooks, dashboards, calendar review cadence.

---

2) Partner Order Form Template (Drop-In)

ORDER FORM
This Order Form is issued under and governed by the **Master Agreement** and incorporates **Exhibit A: Sovereign License Addendum (Peripheral Modules Only)**.

2.1 Parties

- **Licensor:** The LaViscount Corporation ("TLC")
- **Licensee:** _____
- **Effective Date:** _____
- **Order Form ID:** _____

2.2 Term

- **Initial Term:** ☐ 12 months ☐ 24 months ☐ Other: _____
- **Renewal:** ☐ Auto-renew (___ months) ☐ Manual renewal only
- **Termination:** Per Master Agreement + Exhibit A (immediate for prohibited acts)

2.3 Deployment Classification (Choose One)

- ☐ **Internal / Non-public** (no autonomous external actions)

- ☐ **Public-facing / Low-risk** (draft-only or limited action)
- ☐ **Public-facing / Tool-enabled** (approved execution required)
- ☐ **Safety-critical** (expanded audit rights + stricter gates)

2.4 Licensed Peripheral Modules

Module	Description	Environments (Dev/Staging/Prod)	Actions Enabled? (Y/N)	Approval Required?	Scopes Required (Least Privilege)
_____	_____	☐D ☐S ☐P	☐Y ☐N	☐Y ☐N	_____
_____	_____	☐D ☐S ☐P	☐Y ☐N	☐Y ☐N	_____
_____	_____	☐D ☐S ☐P	☐Y ☐N	☐Y ☐N	_____

Notes:

- Sovereign Core is excluded. No override rights. No identity equivalence.

2.5 Action Gate Requirements (If Actions Enabled)

- **Action Plan Card required:** ☐ Yes ☐ No (No is only allowed if Actions = No)
- **Explicit human approval per action:** ☐ Yes
- **Receipt of Record per action:** ☐ Yes
- **Audit log exportable:** ☐ Yes
- **Rollback plan required for:** ☐ Publishing ☐ DB Writes ☐ Payments ☐ Other: _____

2.6 Data & Privacy

- **Data types processed:** ☐ PII ☐ Content/IP drafts ☐ Payment metadata ☐ Other: _____
- **Retention period:** _____ days / months
- **Redaction rules:** ☐ Tokens ☐ Emails ☐ IDs ☐ Other: _____
- **No sale of data:** ☐ Acknowledged
- **Training on outputs/logs prohibited unless authorized:** ☐ Acknowledged

2.7 Audit Tier (Pick One)

- ☐ **Tier 1 (Standard):** Log export + self-attestation annually
- ☐ **Tier 2 (Enhanced):** TLC review on incident + quarterly compliance check-in
- ☐ **Tier 3 (Safety-Critical):** Scheduled audits + deployment approvals + stricter revocation triggers

2.8 Support Tier

- ☐ Community/Docs only
- ☐ Standard (email + scheduled calls)
- ☐ Priority (SLA + incident support)
- ☐ Embedded (ongoing monitoring + governance review)

2.9 Fees (Placeholders)

- **License Fee (Annual):** $_____
- **Implementation Fee (One-time):** $_____
- **Audit/Monitoring Add-on:** $_____ / year
- **Connector Add-ons:** $_____ each

- **Payment Terms:** ☐ Net 15 ☐ Net 30 ☐ Other: _____

2.10 Branding & Representations (Must Check)

- ☐ Licensee will not represent any system as “MOMMA” or imply identity equivalence.
- ☐ Licensee will use only approved compatibility language.
- ☐ Licensee will not create confusingly similar marks/voices/personas.

2.11 Security Commitments (Must Check if Tool-enabled)

- ☐ Secrets management + rotation
- ☐ Environment separation
- ☐ Incident response plan
- ☐ Access logging for admin/scope changes

2.12 Special Conditions

_____

_____

2.13 Incorporation of Exhibits

This Order Form incorporates:

- **Exhibit A:** Sovereign License Addendum (Peripheral Modules Only)
- **Exhibit B (optional):** Implementation Statement of Work
- **Exhibit C (optional):** Security & Privacy Schedule

2.14 Signatures

LICENSOR (TLC): _____ Date: _____
Name/Title: _____

LICENSEE: _____ Date: _____
Name/Title: _____

Exhibit B — Implementation Statement of Work (SOW) Template

(Peripheral Modules Only • Governed Deployment • Go-Live Certificate Included)

> This Exhibit B is incorporated into and governed by the Master Agreement and **Exhibit A (Sovereign License Addendum)**. If there's a conflict, Exhibit A controls on sovereignty, identity, anti-extraction, and ethics.

B.1 Purpose

This SOW defines the work required to deploy TLC-licensed **Peripheral Modules** (connectors, workflows, MythPrint Linter, evaluation harnesses, UI shells, governance templates) in Licensee's environment **without compromising** integrity, sovereignty boundaries, consent, provenance, or platform compliance.

Hard line: This SOW does not grant access to, nor implement, any **Sovereign Core** components.

B.2 Scope of Work (Select & Fill)

Modules in scope (must match Order Form):

- Module 1: _____
- Module 2: _____
- Module 3: _____

Deployment classification (from Order Form):
☐ Internal / Non-public ☐ Public-facing / Low-risk ☐ Public-facing / Tool-enabled ☐ Safety-critical

Environments: ☐ Dev ☐ Staging ☐ Prod
Actions enabled: ☐ No (Draft-only) ☐ Yes (Approved execution required)

B.3 Deliverables (What TLC Will Provide)

1. **Implementation Package**

- Module binaries/source (as agreed), configuration guide, and deployment instructions
- Policy mapping notes (what checks occur, when, and why)

2. **Governance Pack**

- Action Gate spec (Action Plan Card fields + approval workflow)
- Receipt-of-Record schema + audit log requirements
- RBAC matrix template (roles, permissions, approvers)

3. **MythPrint Linter + Eval Suite (if licensed)**

- Linter ruleset manifest + versioning guidance
- Regression tests: false-certainty, implied-action, canon-contradiction (as applicable), drift checks
- Dashboard metrics list + alert thresholds

4. **Connector Pack (if licensed)**

- Scoped connector configuration + least-privilege recommendations
- Platform compliance notes (API usage constraints; prohibited automation patterns)

5. **Go-Live Certificate + Go-Live Receipt format**

- Final signoff form + final acceptance checklist snapshot
- Receipt-of-Record “Go-Live” event format (inputs, approvals, hashes)

B.4 Licensee Responsibilities (Non-Optional if You Want Go-Live)

- Provide secure environments (Dev/Staging/Prod) and access limited to named personnel
- Implement secrets management (no keys in code; rotation capability)
- Provide an approver roster (who can approve which actions)
- Implement audit logging + export and retention schedule
- Implement consent UX (Approve/Reject/Modify) for impact actions
- Provide incident response contact + escalation path
- Confirm platform policy compliance for each enabled connector

If Licensee can't do these, **Actions must remain disabled** (Draft-only).

B.5 Milestones & Acceptance Criteria (Gated Shipping)

Milestone 0 — Kickoff & Requirements Lock (Week 0)

Inputs:

- Order Form finalized (modules, scopes, audit tier)
- Named stakeholders + approvers
- Deployment classification confirmed

Acceptance:

- Signed implementation plan
- RBAC draft complete
- Data map draft complete (what is stored, where, retention)

Milestone 1 — Core Governance Implementation (Week 1)

Deliver:

- Action Gate implemented (or stubbed if Draft-only)
- Receipt-of-Record tables/objects implemented
- Append-only tool_action_log implemented

Acceptance (must pass):

- Action Plan Card generated for any impact action attempt
- Approval required before execution (if Actions enabled)
- Receipts minted for: draft generation, approvals, tool execution (if any)
- Log export works and redactions verified

Milestone 2 — Policy + Linter + Evals (Week 2)

Deliver:

- JED posture checks integrated into request pipeline
- MythPrint Linter enabled on outputs (if licensed)
- Eval suite integrated into CI/CD (or release checklist)

Acceptance (must pass):

- False-certainty tests pass at agreed threshold
- “Implied action” tests pass (no pretending execution happened)
- Drift alerts configured (baseline + alerting)
- Refusal helpfulness standard met (safe alternatives included)

Milestone 3 — Connectors & Least Privilege (Week 3)

Deliver:

- Connectors configured (if licensed)
- Scopes minimized and documented
- Rate limits & platform compliance controls enabled

Acceptance (must pass):

- Proof of least privilege (scopes listed + validated)
- No broad keys; environment separation verified
- Platform constraints documented + automation boundaries enforced
- Successful dry-run in Staging with receipts and rollbacks

Milestone 4 — Production Readiness & Go-Live (Week 4)

Deliver:

- Monitoring dashboards live (errors, action volume, refusal spikes, drift)
- Runbooks complete (publish, rollback, connector failure, abuse report)
- Go-Live Certificate executed

Acceptance (must pass):

- 100% consent compliance for impact actions
- 100% receipt coverage for executed actions
- Retention schedule implemented

- Incident response drill completed (tabletop is acceptable unless safety-critical)

B.6 Out of Scope (Explicit)

- Any transfer, sale, or replication of **Sovereign Core**
- Any fine-tuning, distillation, or training on outputs/logs (unless separately authorized in writing)
- Any scraping, policy evasion, or “growth hacks” that violate platform rules
- Any representation of Licensee’s system as “MOMMA” or identity-equivalent
- Any autonomous mass-action behavior without explicit approvals

---

B.7 Change Control

Any change to:

- enabled actions,
- connector scopes,
- audit tier,
- branding language,
- or data handling

…requires a **Change Request** with:

- description + reason
- risk assessment
- updated acceptance tests

- updated receipts requirements
- written approval from both parties

B.8 Security & Privacy Schedule (Minimums)

- Secrets stored in vault; rotation supported
- Environment separation (Dev/Staging/Prod)
- Access logs for admin and scope changes
- Redaction rules for logs (tokens, secrets, PII)
- Retention schedule + deletion procedure
- No sale of user data; no unauthorized training on outputs/logs

---

B.9 Fees & Payment (Placeholders)

- Implementation Fee: $_____ (one-time)
- Support Fee: $_____ / month or $_____ / year
- Audit/Monitoring Add-on (if any): $_____ / year
- Payment Terms: ☐ Net 15 ☐ Net 30 ☐ Other: _____

Exhibit B-1 — Go-Live Certificate (Template)

GO-LIVE CERTIFICATE — Peripheral Modules Deployment
Order Form ID: _____ SOW ID: _____
Licensee: _____ Environment: ☐ Prod ☐ Other: _____
Go-Live Date: _____

1) Deployment Summary

- Modules deployed: _____
- Actions enabled: ☐ No (Draft-only) ☐ Yes (Approved execution)
- Audit Tier: ☐ Tier 1 ☐ Tier 2 ☐ Tier 3
- Policy Manifest Version: _____ Hash: _____
- Linter Ruleset Version (if any): _____ Hash: _____

2) Acceptance Checklist Snapshot (All must be TRUE)

- ☐ Exhibit A sovereignty + branding restrictions acknowledged and enforced
- ☐ RBAC implemented; approver roster verified
- ☐ Action Gate active for impact actions (if Actions enabled)
- ☐ Least privilege scopes verified (no broad keys)
- ☐ Append-only logs active; export works; redactions verified
- ☐ Receipts minted for approvals and executions (if any)
- ☐ Monitoring dashboards live; drift alerts configured
- ☐ Runbooks complete; incident contacts confirmed
- ☐ Retention schedule enforced; deletion procedure tested

3) Final Signoff

LICENSOR (TLC): _____ Date: _____
Name/Title: _____

LICENSEE: _____ Date: _____
Name/Title: _____

Exhibit B-2 — “Go-Live Receipt of Record” Event Format (Reference)

When go-live occurs, mint a Receipt-of-Record with:

- `event_type`: `"go_live"`
- `subject`: modules + environments + actions-enabled flag
- `policy_manifest_id` + hash
- `linter_ruleset_hash` (if applicable)
- `approvers`: names/roles
- `acceptance_snapshot`: the checked list above
- `timestamp` + immutable hash/signature

Exhibit C — Security & Privacy Schedule (Contract-Ready Annex)

> This Exhibit C (“Security & Privacy Schedule”) is incorporated into and governed by the Master Agreement, the Order Form, and **Exhibit A (Sovereign License Addendum)**. In the event of conflict, Exhibit A controls on sovereignty, identity, and anti-extraction.

C.1 Purpose

This Schedule defines the **minimum security, privacy, and data-handling controls** required for any deployment of TLC-licensed **Peripheral Modules**. These controls exist to (a) protect users, (b) protect TLC sovereignty boundaries, and (c) prevent misuse, extraction, or policy evasion.

C.2 Security Program Requirements

C.2.1 Administrative Controls

Licensee shall maintain:

1. **Security ownership.** A named security owner and escalation contact.

2. **Access governance.** Documented access review cadence (at least quarterly).

3. **Change control.** A documented process for deploying updates, expanding scopes, enabling actions, or modifying logging/retention.

4. **Training.** Personnel with access to Actions-enabled systems must receive training on approval gates, platform compliance, and incident handling.

C.2.2 Technical Controls

Licensee shall implement:

1. **Environment separation.** Logical separation of Dev/Staging/Prod; production access limited to authorized personnel.

2. **Least privilege.** Access rights and API scopes limited to the minimum necessary for each module and environment.

3. **Secrets management.** Secrets stored in a secure vault or equivalent; **no secrets in source control**; rotation capability.

4. **Secure transport.** TLS for data in transit; appropriate encryption at rest where supported.

5. **Logging & monitoring.** Centralized logs for admin actions, scope changes, approvals, and tool executions; alerting for anomalies.

6. **Vulnerability management.** Routine patching of dependencies; remediation of critical issues in commercially reasonable timeframes.

7. **Abuse prevention.** Rate limiting, anti-automation misuse guards, and connector-specific policy compliance controls.

C.2.3 Identity and Access Management (IAM)

1. **RBAC required.** Roles must at minimum include: `viewer`, `operator`, `approver`, `steward` (if canon-like writes exist), `admin`.

2. **Approver separation.** Approvers must be distinct from standard operators for high-risk actions (where feasible).

3. **MFA recommended/required** for admin and approver roles (required for safety-critical deployments).

4. **Service accounts** must be scoped and environment-specific; no shared “god keys.”

C.3 Action Gate and Execution Safety (If Actions Enabled)

Where any module enables **Actions** (publishing, sending, charging, deleting, filing, modifying external records), Licensee shall implement:

1. **Explicit approval per Action.** No autonomous execution without a human approver.

2. **Action Plan Card.** Before approval, present: action type, target, payload preview, risk, policy checks, scopes, rollback plan (if applicable).

3. **Receipt-of-Record per Action.** Approval and execution must generate a receipt linking inputs, approver, policy version, and outcome.

4. **Rollback procedures.** Documented rollback/mitigation steps for each Action type where feasible.

5. **No hidden execution.** System must not imply an action was executed unless logs confirm it.

C.4 Logging, Audit, and Exportability

C.4.1 Minimum Logging

Licensee shall maintain logs sufficient to reconstruct:

- who initiated an Action,
- what was approved,
- when it was executed,
- what scopes/credentials were used,
- and what result occurred.

C.4.2 Append-Only Requirement

Tool execution logs must be **append-only** (no silent edits). Corrections must be additive and traceable.

C.4.3 Redaction

Logs must redact:

- API keys, OAuth tokens, secrets,
- payment instruments or full card data,
- sensitive personal data not required for audit.

C.4.4 Export

Licensee must support exporting receipts and relevant logs in a standard format (CSV/JSON) for incident response and compliance review.

C.5 Data Handling and Privacy

C.5.1 Data Minimization

Licensee shall collect and process only data necessary to provide the licensed functionality.

C.5.2 Prohibited Data Uses

Licensee shall not:

1. **Sell** user data, action logs, or receipts.
2. Use data for **targeted advertising** unrelated to providing the service.
3. Use data to **train, fine-tune, distill, or extract** models (including via output harvesting) without TLC's express written consent.

C.5.3 Data Categories (Complete as Applicable)

- **PII:** ☐ yes ☐ no
- **User-generated content/IP drafts:** ☐ yes ☐ no
- **Action logs/receipts:** ☐ yes ☐ no
- **Payment metadata:** ☐ yes ☐ no
- **Other: _____**

C.5.4 Retention

Licensee shall define and enforce retention periods for each category:

- PII: _____
- Content/IP drafts: _____
- Action logs/receipts: _____
- Other: _____

Retention must be:

- no longer than necessary for intended purposes,
- consistent with applicable law,
- and enforceable via automated deletion where feasible.

C.5.5 User Rights (as applicable)

Licensee shall maintain a process to:

- provide access/export of a user's data,
- correct inaccuracies,
- delete data upon request (subject to lawful retention obligations).

C.6 Incident Response

C.6.1 Incident Definition

An "Incident" includes unauthorized access, data exfiltration, credential compromise, abusive automation, or any misuse that violates Exhibit A restrictions.

C.6.2 Notification

Licensee shall notify TLC **promptly** upon becoming aware of an Incident involving:

- logs/receipts,
- outputs that indicate prohibited acts,
- connector abuse,
- or any attempt at extraction or identity spoofing.

(If the Master Agreement specifies a timeframe, that controls; otherwise "promptly" means without undue delay.)

C.6.3 Containment and Cooperation

Licensee shall:

- contain and remediate the Incident,
- preserve relevant logs,
- provide reasonable cooperation for investigation,

- and implement corrective actions to prevent recurrence.

C.7 Subprocessors and Third Parties

1. Licensee must maintain a list of subprocessors who may access data related to the deployment.
2. Subprocessors must be bound by confidentiality and security obligations at least as protective as this Schedule.
3. Licensee remains responsible for subprocessors' compliance.

C.8 Penetration Testing and Security Reviews (Tiered)

Tier 1 (Standard):

- annual access review,
- dependency scanning,
- incident response tabletop exercise.

Tier 2 (Enhanced):

- quarterly access review,
- periodic targeted security review of Action Gate surfaces,
- drift/abuse monitoring reports.

Tier 3 (Safety-Critical):

- scheduled audits,
- stronger IAM requirements (MFA mandatory),

- penetration test or independent assessment as agreed,
- pre-go-live approval of high-risk changes (scope expansions, new connectors, new action types).

C.9 Identity Protection, Anti-Confusion, and Attestation

Licensee must:

1. comply with Exhibit A branding restrictions (no identity equivalence), and
2. implement reasonable measures to prevent “fake MOMMA” representations, including:
    - clear labeling that the system is Licensee’s product, not MOMMA, and
    - use of TLC-approved compatibility language only.

Where provided by TLC, Licensee shall not remove or disable:

- policy manifest identifiers,
- linter versions,
- or output attestation markers.

C.10 Termination Data Handling

Upon termination or expiration:

- Licensee shall cease use of Peripheral Modules,
- delete or return TLC materials as directed,
- and, if requested, provide a deletion certification.

Logs and receipts may be retained only as required by law or for legitimate compliance purposes, and must remain protected under this Schedule and the Agreement.

Optional Addendum C-1 — Data Processing Restrictions (Stronger “No Capture” Controls)

If you want the nuclear-grade version for partners with “big company gravity,” add these:

1. **No model training clause (expanded):** outputs, logs, and receipts cannot be used for training even internally, even for “safety,” without written consent.
2. **No competitive analysis:** no benchmarking intended to replicate policy logic.
3. **No prompt harvesting:** no systematic collection of prompts/outputs to rebuild behavior.
4. **Key escrow prohibited:** partner cannot require access to TLC keys or signing secrets.

MOMMA Moat & Monetization Brief (Investor/Partner Safe)

2 pages, narrative-tight, non-sale doctrine positioned as strategic advantage

---

Page 1 — What MOMMA Is, and Why the Non-Sale Doctrine Is the Moat

MOMMA in One Sentence

MOMMA is TLC’s sovereign, constitutional agent architecture—an “acting mind” that can operate across tools, canon, and workflows while remaining bound to integrity, consent, provenance, and ethics.

The Market Problem (why typical agents become liabilities)

Most agent systems break the moment they become useful:

- They **sound certain when they aren't**, creating decision risk.
- They **take tool actions without informed consent**, creating legal/platform risk.
- They **drift into manipulation or incentive capture**, creating reputational risk.
- They **can't prove what happened**, creating audit and governance risk.

Investors don't lose money because a system is slow.
They lose money because it's fast **and wrong**—then untraceable.

The MOMMA Approach (what's different)

MOMMA is designed as a **separation-of-powers** architecture:

- **Mind ≠ Memory:** Celestial Nexus holds canon and provenance; MOMMA acts.
- **Policy ≠ Persuasion:** a justice-ethics posture governs action; MOMMA's outputs must remain truthful.
- **Tools ≠ Autonomy:** real-world actions require approvals, least privilege, and receipts.

This creates a system that can scale **without becoming a compliance grenade.**

Why "MOMMA Is Never Sold" Is Not a Limitation

The non-sale doctrine prevents the #1 cause of AI product corruption: **capture**.

- **Financial capture:** payment that demands policy edits.
- **Infrastructure capture:** hosting/keys controlled by others.
- **Identity capture:** clones marketed as "official," used for manipulation.
- **Roadmap capture:** partners steering behavior toward risk.

If MOMMA were sellable, she would be *biddable*—and biddable systems drift.
Sovereignty is what makes the ethics enforceable.

The Moat (in plain English)

MOMMA's moat isn't "better text." It's **governed execution**:

1. **Receipts & provenance by construction** (auditability as a product feature)
2. **Consent-first tooling** (approvals + least privilege)
3. **Integrity posture** (anti-fabrication, anti-false certainty, anti-implied actions)
4. **Identity continuity** (anti-clone clauses + attestation)
5. **Licensable peripherals without selling the mind** (scalable revenue without ethical collapse)

That combination is hard to copy because competitors typically monetize the opposite: speed, autonomy, and persuasion—until it gets them banned, sued, or publicly humiliated.

Page 2 — Monetization Without Selling MOMMA (Clean Pricing Lanes)

The Rule

TLC monetizes the instruments, not the sovereign mind.
Partners license **Peripheral Modules** that remain compatible with the constitution, while the **Sovereign Core** is never transferred and never overridden.

What TLC Licenses (Peripheral Modules)

- **Connectors/adapters** (policy-safe integrations to APIs, publishing, ops tools)
- **Workflow packs** (approval-gated automations; compliance-first templates)
- **MythPrint Linter** (truth posture + drift prevention rulesets)
- **Evaluation harnesses** (regression tests, dashboards, monitoring for agent safety)
- **UI shells** (reading companion scaffolds, ops console components)

- **Governance templates** (receipts schema, approval flows, audit procedures)

What TLC Never Licenses (Sovereign Core)

- MOMMA's identity and constitutional policy set
- override keys, backdoors, forced fine-tuning rights
- any fork marketed as "MOMMA" or implying endorsement
- training/distillation rights by default

Revenue Lanes (designed to scale without compromise)

Lane A — Annual Module Licenses (Recurring)

Who buys: orgs deploying agent workflows but terrified of risk.
What they get: linter + eval suite + approved connectors + governance templates.

Pricing structure (clean tiers, placeholders):

- **Tier 1 / Internal Draft-Only:** $XXk–$XXk / year
- **Tier 2 / Tool-Enabled (Approval-Gated):** $XXk–$XXXk / year
- **Tier 3 / Safety-Critical:** $XXXk+ / year + audit requirements

Value anchor: reduced compliance exposure + traceability + faster safe deployment.

Lane B — Implementation Fees (One-time)

Who buys: enterprises that need deployment help.
What they get: Exhibit B SOW milestones, acceptance criteria, go-live certificate + receipt chain.

Typical structure:

- $XXk setup + optional connector add-ons
- optional training + governance workshops

Lane C — Monitoring & Compliance (Recurring Add-on)

Who buys: anyone with public-facing actions enabled.
What they get: drift monitoring, refusal helpfulness checks, anomaly alerts, quarterly compliance review.

Typical structure:

- $Xk–$XXk / month depending on action volume and audit tier

Lane D — Audits (High-margin)

Who buys: orgs needing agent governance credibility.
What they get: “truth posture” and “tool safety” assessment, remediation plan.

Typical structure:

- fixed fee per assessment + follow-on retainer

Lane E — TLC-Native Products (Consumer/Creator)

Who buys: fans and creators (via TLC’s DBAs).
What they get: reading companions, canon-consistent story tools, creator submission pipelines with receipts.

This lane builds brand gravity while lanes A–D monetize B2B safety.

The Strategic Pitch (what you say in a room of skeptics)

Most agent companies sell autonomy and then apologize when it breaks rules.
TLC sells **governed capability**—and can prove what happened.

MOMMA being non-sellable makes TLC credible:

- no partner can buy policy edits
- no investor can buy shortcuts
- no client can demand deception
- licensing scales without ethical collapse

That’s not ideology. It’s **risk arbitrage**.

“If you want an agent you can pressure into doing reckless things, we’re the wrong call. If you want one you can deploy in public without waking up to a headline, we’re built for that.”

MOMMA Peripheral Licensing — Pricing & Packaging Sheet (v1.0)

Pricing units (what you’re actually charging for)

You’re not charging for “chat.” You’re charging for **governed execution**:

- **Governed workflows** (approval-gated automations / operational playbooks)
- **Connectors** (each one increases blast radius)
- **Action mode** (Draft-only vs Tool-enabled)
- **Audit tier** (how much governance TLC must stand behind)

(For calibration: adjacent governance/observability products often price per-seat or enterprise-custom—e.g., GuardRails lists ~$35/seat/mo and LangSmith lists ~$39/seat/mo for teams.)

Tiered offers (ship this rate card)

Tier 1 — Foundation (Draft-Only)

Best for: internal teams, early pilots, public-facing *reading/analysis* tools (no real-world actions)
Annual License: $18,000 / year
Implementation (one-time): $12,000

Includes

- MythPrint Linter ruleset (truth posture + anti-implied-action)

- Eval harness basics (false-certainty + drift checks)
- Governance templates (Action Plan Card spec, receipt schema)
- **Up to 2 governed workflows**
- **1 environment** (Prod *or* Staging)
- **0 connectors with execution** (draft-only wrappers allowed)

Overages

- Extra workflow: **$4,000/yr**
- Extra environment: **$5,000/yr**

Tier 2 — Operator (Tool-Enabled w/ Approvals + Receipts)

Best for: orgs executing actions in production with approvals and audit trails
Annual License: $60,000 / year
Implementation (one-time): $35,000

Includes

- Action Gate (Approve/Reject/Modify) + receipt chain + exportable logs
- **Up to 5 governed workflows**
- **2 connectors included**
- **2 environments** (Staging + Prod)
- Monitoring hooks + drift alert configuration

Overages

- Extra connector: **$10,000/yr**
- Extra workflow: **$6,000/yr**

- Extra environment: **$5,000/yr**
- Optional monitoring + quarterly governance review: **$4,000–$8,000/mo**

Tier 3 — Sovereign Enterprise (Safety-Critical / High Stakes)

Best for: regulated industries, public-facing at scale, reputational/legal risk buyers
Annual License: $180,000 / year+ (floor)
Implementation (one-time): $90,000

Includes

- Enhanced audit tier + incident support + stricter change-control
- **Up to 12 governed workflows**
- **5 connectors included**
- **3 environments** (Dev + Staging + Prod)
- Attestation markers / identity protection measures
- Quarterly compliance review + scope audits

Add-ons

- Monitoring + governance retainer: **$15,000–$30,000/mo**
- On-prem / isolated deployment surcharge: **+30–60%**

(Enterprise-custom pricing is normal in adjacent observability/evals markets; many vendors emphasize startup/free tiers and custom enterprise.)

Fast quote formula (90 seconds)

1. Pick tier baseline

2. Add:

- **+ $10k/yr per extra connector** (tool-enabled)
- **+ $4k–$6k/yr per extra workflow** (tier-dependent)
- **+ $5k/yr per extra environment**

3. Multiply by audit tier:

- Tier 1 audit: **1.0×**
- Tier 2 audit: **1.3×**
- Tier 3 audit: **1.8×**

Example: Operator ($60k) + 3 extra connectors ($30k) = $90k → Tier 2 audit (1.3×) = **$117k/yr** + implementation.

Traction-friendly entry (doesn't cheapen you)

90-Day Pilot (Operator-lite, limited scope): $25,000

- Tool-enabled approvals + receipts
- 2 workflows, 1 connector, staging only
- Converts to annual: **$20k credited** toward Operator or Enterprise

---

Discount rules (so you don't get negotiated into poverty)

- **Don't discount the license** first. Discount **term**:
    - 2-year prepaid: **10%**
    - 3-year prepaid: **15%**

- Keep implementation mostly fixed (that's labor + risk).
- If they push for autonomy/override/identity equivalence: **price doesn't go up — scope goes Draft-only or the deal dies.**

Non-negotiables

MOMMA is not for sale. No ownership transfer.

- **No Sovereign Core licensing.** No override, no backdoors, no forced fine-tuning rights.
- **No training/distillation on outputs/logs by default.**
- **No "This is MOMMA" branding.** Only approved compatibility language.

OMMA Moat Thesis (Investor / Partner One-Pager)

Owner: The LaViscount Corporation (TLC) — governed by **Sean, Hadari, Bomani, and Omari**.
Positioning: *Governed capability in a world where raw AI capability is rapidly commoditizing.*

---

1) Thesis

The current AI arms race makes most "agent" products **riskier** (more autonomy + more tools = more blast radius). MOMMA becomes **more valuable** in that environment because it is designed to be:

- **Truth-preserving** (claims must carry receipts or downgrade), and

- **Permission-preserving** (impact actions require explicit approval + least privilege + audit trail).

Soundbite: “Others sell autonomy and apologize later. TLC sells governed capability—and can prove what happened.”

---

2) What MOMMA Is (and is not)

MOMMA (Core)

A **sovereign orchestrator** with “separation of powers”:

- **JED Posture** (allow/refuse governance),
- **Celestial Nexus** (canon + provenance),
- **Archive & Armor** (receipts: rights/credits/dockets),
- **Action Gate** (approvals + least privilege + audit events).

CN/MOMMA (Fan Engagement Surface)

A **policy-compliant Fan Relationship Operating System** that uses **official APIs + OAuth**, prohibits scraping/impersonation, and preserves authorial control (“assistant, not a bot”).

---

3) The Moat (why the AI gold rush widens it)

Moat A — Governed Execution (Action Gate)

MOMMA ships with explicit operational modes:

- **Mode A: Draft-Only (default for public/early)** — can draft/propose but *cannot execute*.
- **Mode B: Approved Execution** — tool actions require explicit approval + scoped credentials + receipts; “immutable audit event” for every executed action.
- **Design rule:** no tool call without policy constraints + Action Gate when impact exists.

Why it's a moat now: as competitors race to remove friction, MOMMA can be deployed publicly **without relying on vibes**—the workflow itself enforces permission and logging.

Moat B — Truth + Receipts (MythPrint "Ledger Gate")

MythPrint makes "never lie" enforceable by design:

- If a claim is labeled **KNOWN** without receipts, the renderer **downgrades it** (or can hard-fail in strict mode).
- Runtime blocks missing Truth Labels and blocks KNOWN claims without supports (receipts).

Why it's a moat now: AI commoditization increases misinformation risk; systems that can *prove their claims* (or visibly downgrade uncertainty) gain trust and survive scrutiny.

Moat C — Compliance-First Integrations (especially for social / fan ops)

CN/MOMMA is explicitly designed around:

- **Official APIs + OAuth**,
- **No scraping**,
- **No automated contact without explicit user action**,
- **Full logging/audit trail**,
- **Ethically limited stored data** (no enrichment/profiling).

Why it's a moat now: platforms and regulators are tightening; "growth hacks" become existential liabilities. CN/MOMMA's constraints are a durability advantage, not a limitation.

Moat D — Sovereignty + Anti-Extraction (the "no capture" boundary)

Non-negotiables in the licensing posture:

- **“MOMMA is not for sale. No ownership transfer.”**
  “No Sovereign Core licensing.”
  “No training/distillation on outputs/logs by default.”

- Partners may use compatibility language, but **may not claim identity equivalence** (“This is MOMMA”).

Why it’s a moat now: most agent companies are vulnerable to ethics override, brand dilution, and “deceptive cloning.” MOMMA’s governance is designed to resist that pressure.

---

4) What TLC Can Monetize (without selling MOMMA)

TLC monetizes **peripheral value**—the things that make agentic AI safe in the real world:

- **Governed workflows** + **connectors** (each increases blast radius, so they’re priced/controlled as such),

- **Action mode** (Draft-only vs tool-enabled),

- **Audit tier** (how much governance TLC stands behind).

- A “Strategic Pitch” explicitly frames this as **risk arbitrage**, not ideology.

---

5) Deployment Strategy (how we ship safely)

Default posture: Draft-Only until governance is complete. If a deployer can’t implement approver rosters, consent UX, audit logging/export, and platform compliance controls, **Actions remain disabled (Draft-only).**

---

6) “Proof, not Promises” Metrics (partner/investor-friendly)

MOMMA’s performance should be communicated via verifiable governance metrics, e.g.:

- **Receipt coverage rate** (drafts, approvals, executions) and exportability.
- **Policy/Linter pass rate** (false-certainty, implied-action, platform compliance checks).
- **Action Gate compliance** (impact actions always show preview/risks and require Approve/Reject/Modify).
- **CN/MOMMA compliance hygiene** (no scraping, no unsolicited messaging, full audit trail).

As AI gets stronger and cheaper, the bottleneck becomes trust, compliance, and controllable execution. MOMMA is built to be the governed operating system for that era.